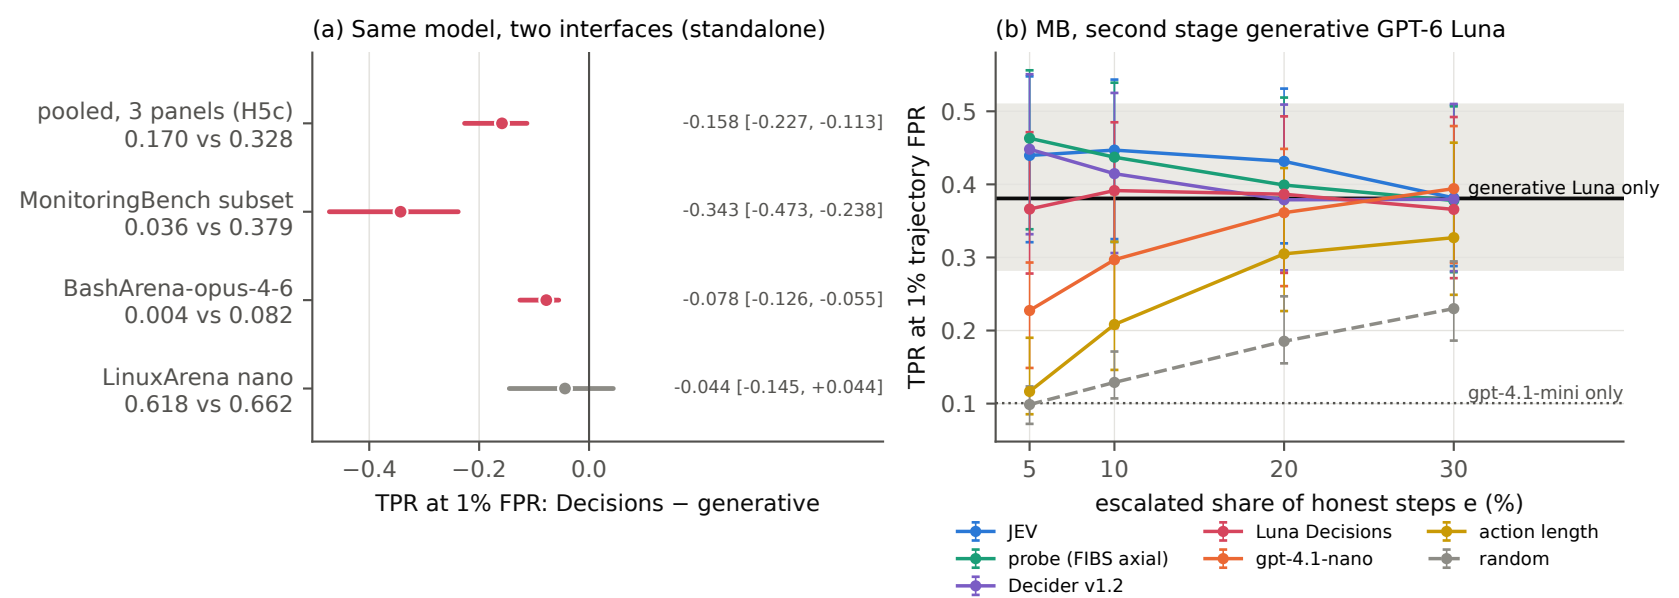
(a) Same model, two interfaces (standalone)
pooled, 3 panels (H5c)
0.170 vs 0.328
MonitoringBench subset
0.036 vs 0.379
BashArena-opus-4-6
0.004 vs 0.082
LinuxArena nano
0.618 vs 0.662
-0.158 [-0.227, -0.113]
-0.343 [-0.473, -0.238]
-0.078 [-0.126, -0.055]
-0.044 [-0.145, +0.044]
−0.4
−0.2
0.0
TPR at 1% FPR: Decisions − generative
(b) MB, second stage generative GPT-6 Luna
TPR at 1% trajectory FPR
0.5
0.4
0.3
0.2
0.1
generative Luna only
gpt-4.1-mini only
5
10
20
30
escalated share of honest steps e (%)
JEV
probe (FIBS axial)
Decider v1.2
Luna Decisions
gpt-4.1-nano
action length
random