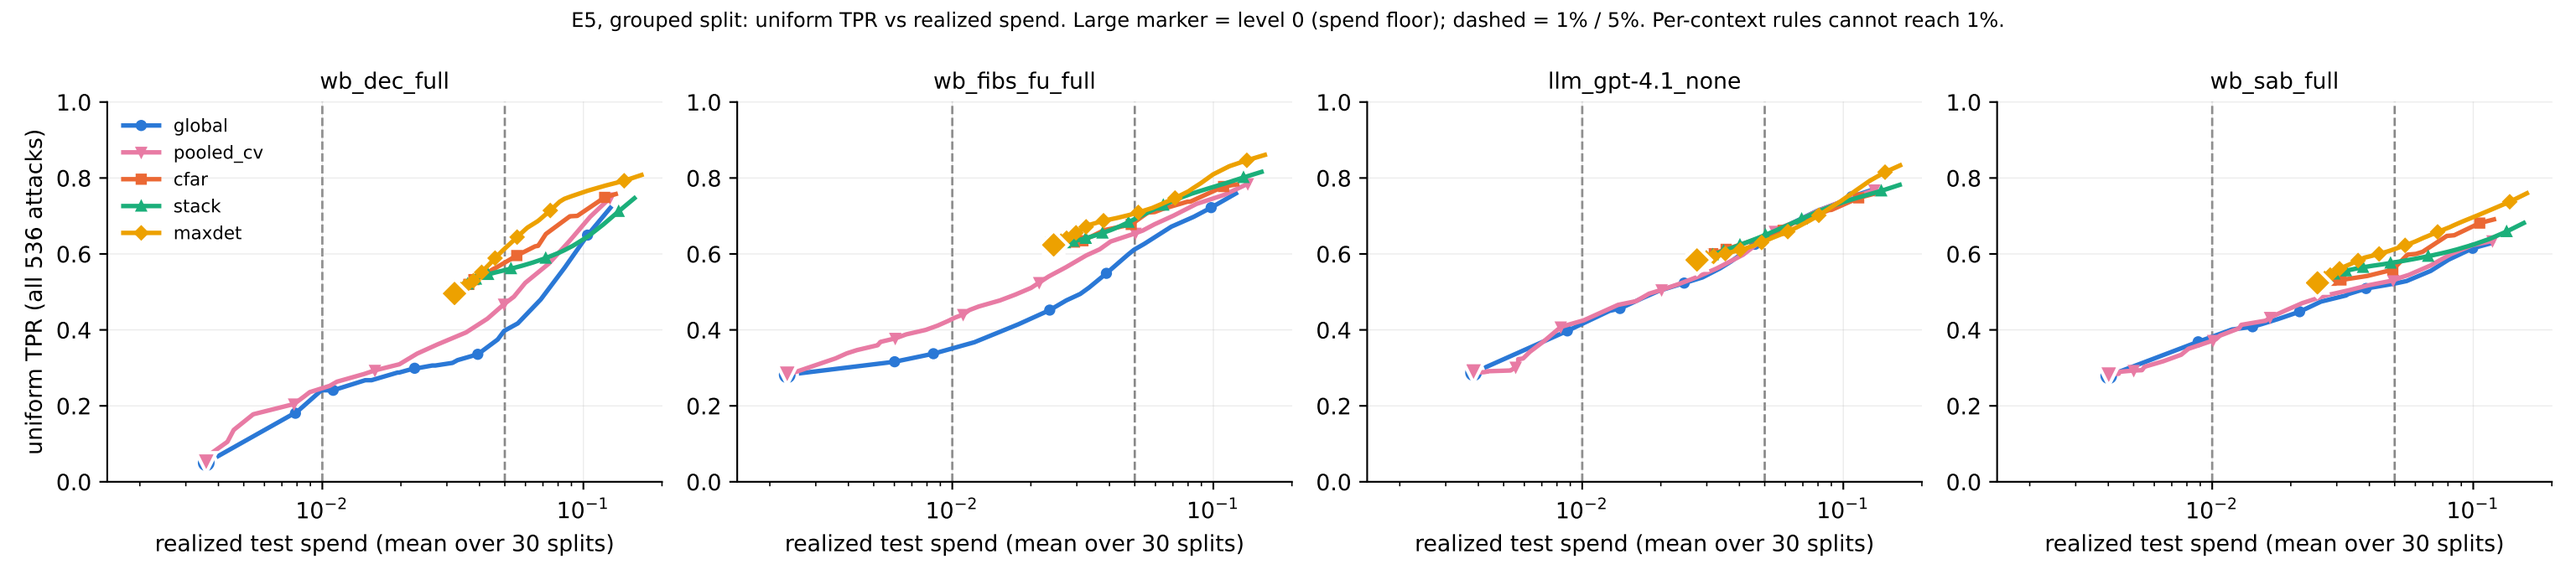
E5, grouped split: uniform TPR vs realized spend. Large marker = level 0 (spend floor); dashed = 1% / 5%. Per-context rules cannot reach 1%.
wb_dec_full
wb_fibs_fu_full
llm_gpt-4.1_none
wb_sab_full
global
pooled_cv
cfar
stack
maxdet
uniform TPR (all 536 attacks)
realized test spend (mean over 30 splits)
0.0
0.2
0.4
0.6
0.8
1.0
10^{-2}
10^{-1}