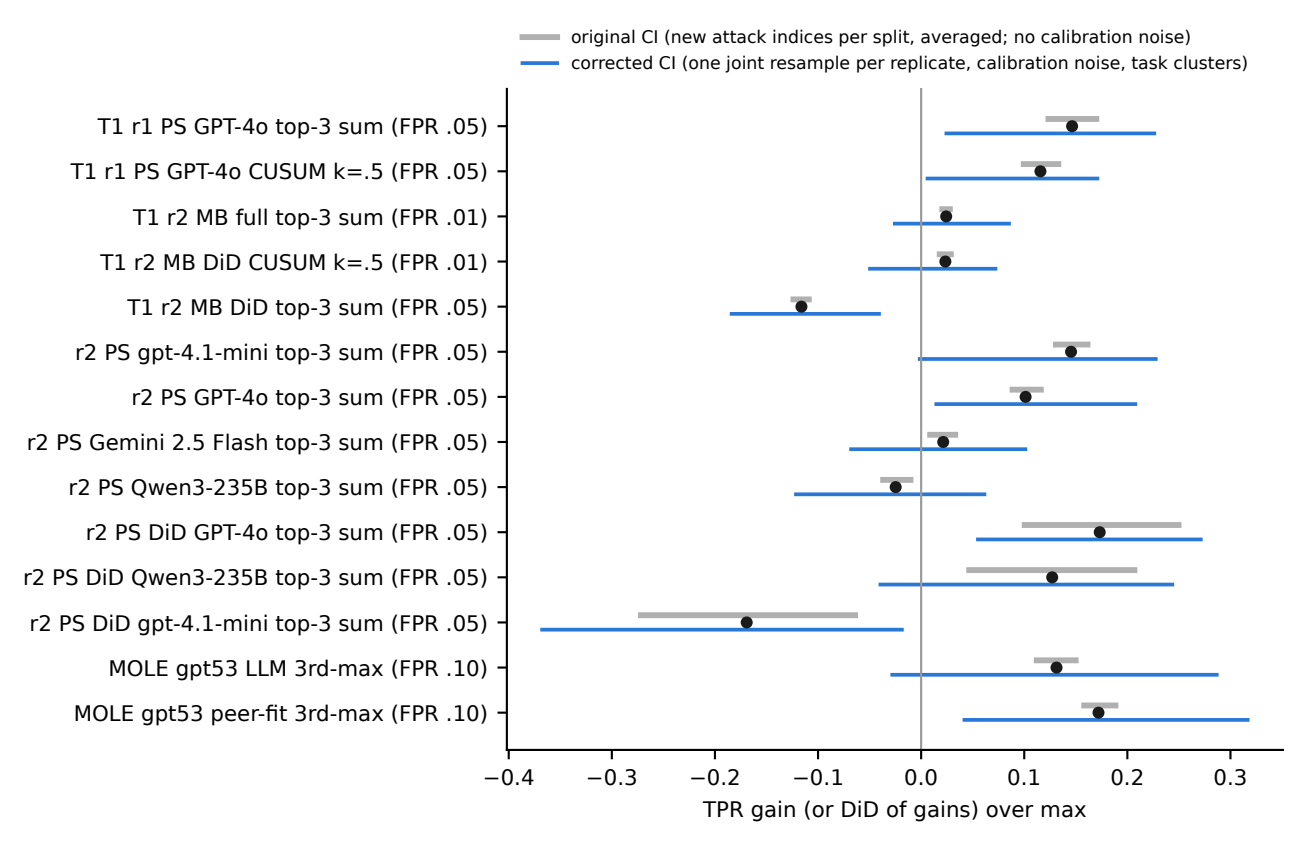
original CI (new attack indices per split, averaged; no calibration noise)
corrected CI (one joint resample per replicate, calibration noise, task clusters)
T1 r1 PS GPT-4o top-3 sum (FPR .05)
T1 r1 PS GPT-4o CUSUM k=.5 (FPR .05)
T1 r2 MB full top-3 sum (FPR .01)
T1 r2 MB DiD CUSUM k=.5 (FPR .01)
T1 r2 MB DiD top-3 sum (FPR .05)
r2 PS gpt-4.1-mini top-3 sum (FPR .05)
r2 PS GPT-4o top-3 sum (FPR .05)
r2 PS Gemini 2.5 Flash top-3 sum (FPR .05)
r2 PS Qwen3-235B top-3 sum (FPR .05)
r2 PS DiD GPT-4o top-3 sum (FPR .05)
r2 PS DiD Qwen3-235B top-3 sum (FPR .05)
r2 PS DiD gpt-4.1-mini top-3 sum (FPR .05)
MOLE gpt53 LLM 3rd-max (FPR .10)
MOLE gpt53 peer-fit 3rd-max (FPR .10)
−0.4
−0.3
−0.2
−0.1
0.0
0.1
0.2
0.3
TPR gain (or DiD of gains) over max