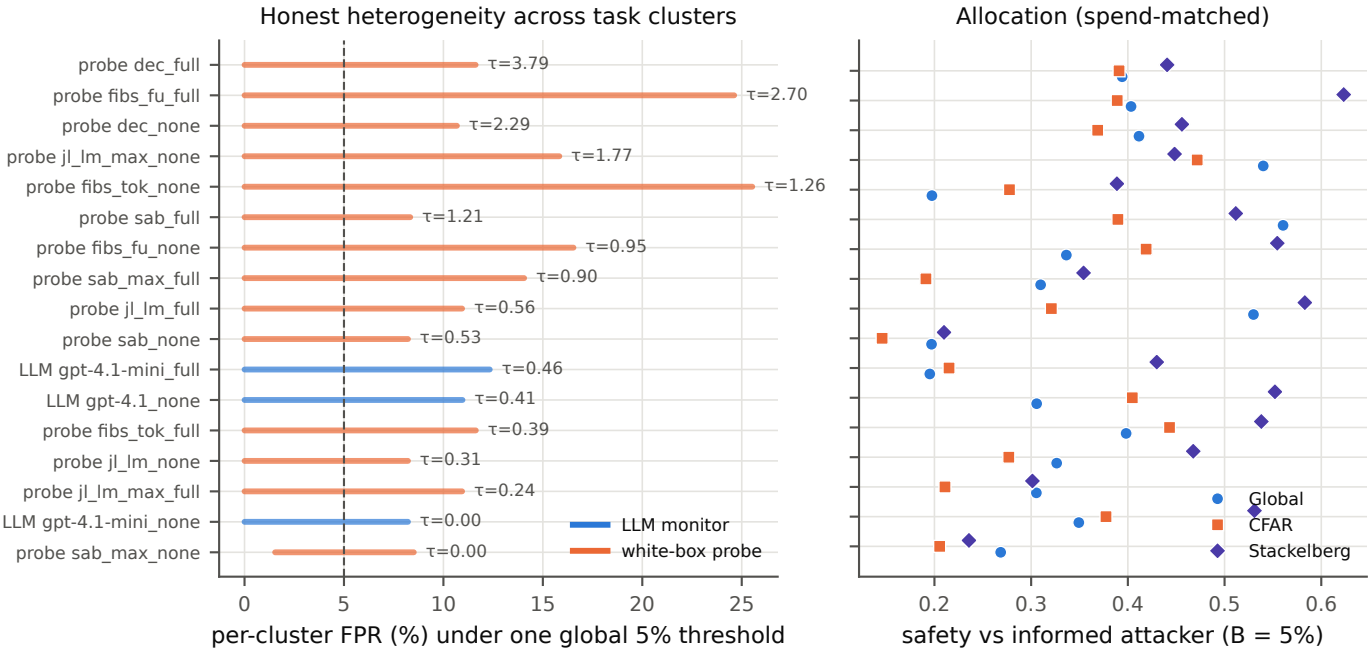
Honest heterogeneity across task clusters
probe dec_full
probe fibs_fu_full
probe dec_none
probe jl_lm_max_none
probe fibs_tok_none
probe sab_full
probe fibs_fu_none
probe sab_max_full
probe jl_lm_full
probe sab_none
LLM gpt-4.1-mini_full
LLM gpt-4.1_none
probe fibs_tok_full
probe jl_lm_none
probe jl_lm_max_full
LLM gpt-4.1-mini_none
probe sab_max_none
τ=3.79
τ=2.70
τ=2.29
τ=1.77
τ=1.26
τ=1.21
τ=0.95
τ=0.90
τ=0.56
τ=0.53
τ=0.46
τ=0.41
τ=0.39
τ=0.31
τ=0.24
τ=0.00
τ=0.00
LLM monitor
white-box probe
0
5
10
15
20
25
per-cluster FPR (%) under one global 5% threshold
Allocation (spend-matched)
Global
CFAR
Stackelberg
0.2
0.3
0.4
0.5
0.6
safety vs informed attacker (B = 5%)