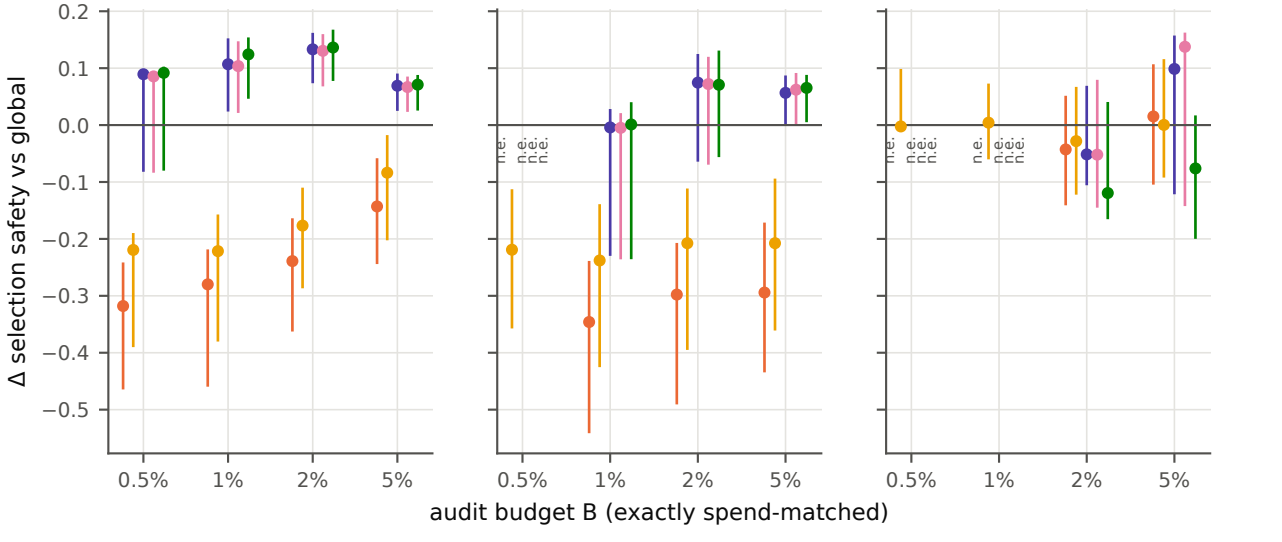
7 settings, Gemini
(328-1,396 honest / context)
ControlArena setting x agent
(~250 honest / context)
public LinuxArena env, gpt-5-nano
(~217 honest / context)
Δ selection safety vs global
0.2
0.1
0.0
−0.1
−0.2
−0.3
−0.4
−0.5
n.e.
0.5%
1%
2%
5%
audit budget B (exactly spend-matched)
CFAR
Partial pooling (EB n0)
Stackelberg, equal damage (m0 = 10)
Stackelberg, unshrunk
Stackelberg, severity-weighted (m0 = 10)