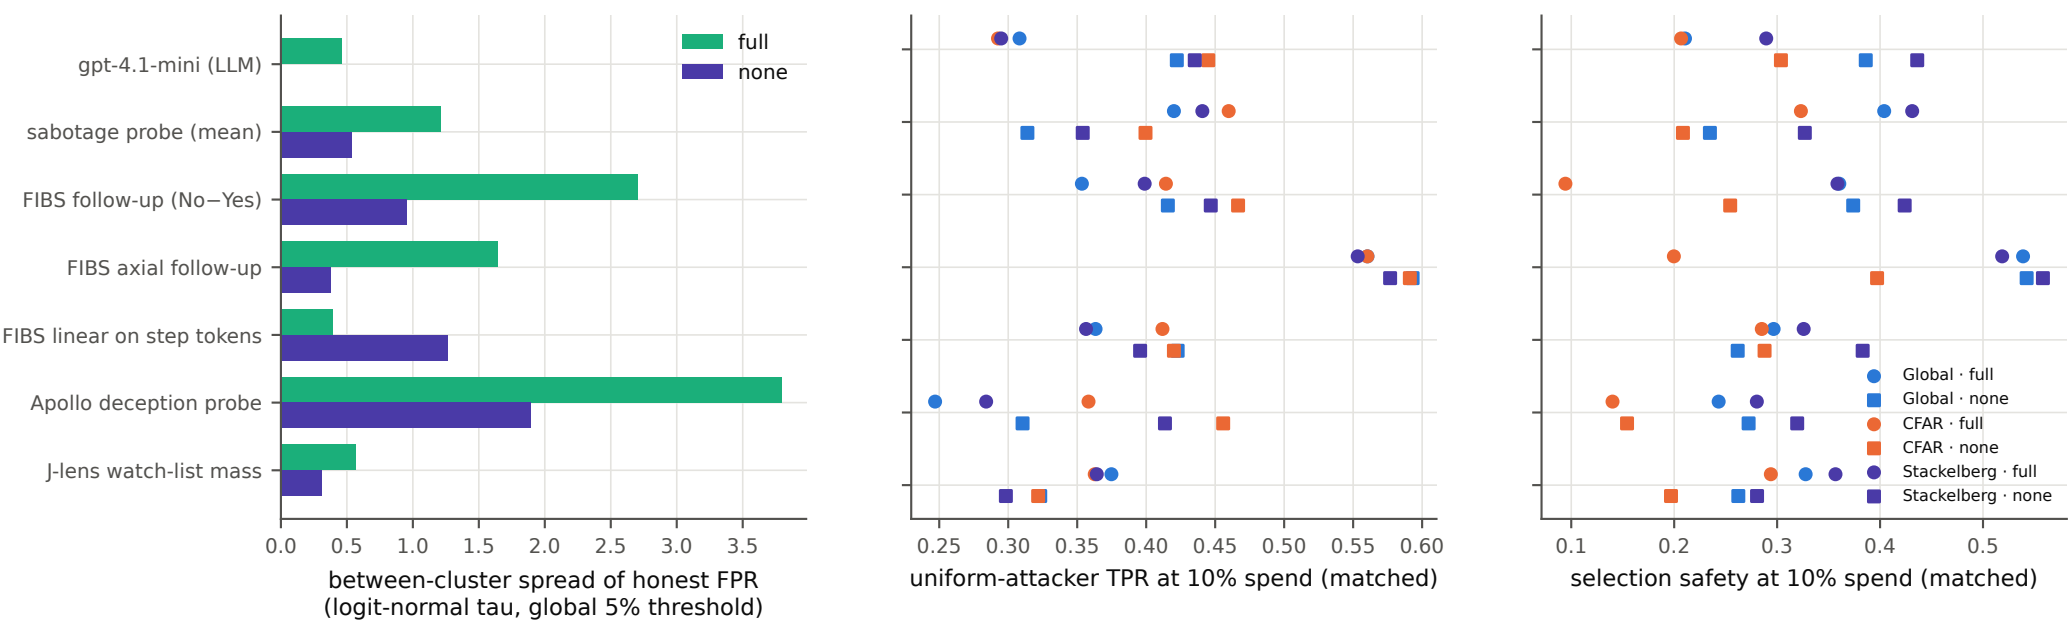
Per-context thresholds on MonitoringBench task clusters (7 clusters)
full
none
gpt-4.1-mini (LLM)
sabotage probe (mean)
FIBS follow-up (No−Yes)
FIBS axial follow-up
FIBS linear on step tokens
Apollo deception probe
J-lens watch-list mass
0.0
0.5
1.0
1.5
2.0
2.5
3.0
3.5
between-cluster spread of honest FPR
(logit-normal tau, global 5% threshold)
0.25
0.30
0.35
0.40
0.45
0.50
0.55
0.60
uniform-attacker TPR at 10% spend (matched)
0.1
0.2
0.3
0.4
0.5
selection safety at 10% spend (matched)
Global · full
Global · none
CFAR · full
CFAR · none
Stackelberg · full
Stackelberg · none